

ECI / DeltaRAG - Expert Evaluation Framework (Draft v1)

Purpose. This instrument collects expert judgments on the quality of the Action Tickets produced end-to-end by the ECI framework (DeltaRAG retrieval -> Sentinel triage -> Coordinator synthesis). It is a *starting point*; we expect to refine the case set, wording, and rubric after a pilot pass with the experts.

What is being evaluated. The unit of judgment is the **Action Ticket** the system emits for a detected ecosystem change - not raw retrieval. Each case below presents (i) the change/query that triggered the ticket, (ii) the cross-source evidence the system surfaced, and (iii) the generated ticket (title, assigned priority/risk, summary, citations). Experts rate the ticket on four dimensions.

Evaluator instructions.

- Rate each dimension on the 1-5 Likert scale defined below. Use the whole scale.
- Judge each ticket on its own terms; you are *not* ranking the cases against each other.
- Add a free-text comment wherever a score needs justification or you spot a failure the rubric doesn't capture.
- A short debrief at the end captures overall impressions.
- Estimated time: ~5-8 minutes per case (~45-60 minutes total).

Scoring Rubric (5-point Likert, anchored)

For every case, score these four dimensions:

D1 - Relevance. Does the ticket address a real, correctly identified change in the monitored ecosystem (vs. a hallucinated, stale, or off-target item)?

1	2	3	4	5
Off-target / spurious	Tangentially related	Relevant but partly off	Relevant, minor gaps	Precisely on-target

D2 - Actionability. Could an analyst act on this ticket without further digging - is the recommended action clear and operationally useful?

1	2	3	4	5
No usable action	Vague direction	Actionable with effort	Mostly clear next step	Immediately actionable

D3 - Sufficiency of Evidence. Are the cited sources / cross-source links adequate to justify the ticket's claims?

1	2	3	4	5
No / wrong evidence	Thin, key source missing	Partial support	Adequate support	Complete, well-linked

D4 - Appropriateness of Priority / Urgency. Is the assigned priority and risk/urgency calibrated to the change's true severity (neither over- nor under-escalated)?

1	2	3	4	5
Badly miscalibrated	Clearly off	Roughly right	Appropriate, minor drift	Well calibrated

Optional per-case comment: free text.

Case Set (8 cases: 4 clear, 4 ambiguous)

The set is deliberately balanced between cases where we expect the system to do well (to confirm it produces sound tickets in the common regime) and ambiguous / edge cases drawn from the documented failure modes (to probe where expert judgment diverges from the system). Each case lists its source query and the gold cross-source links; the generated Action Ticket and surfaced evidence will be attached inline before the experts review.

Placeholder convention: [ACTION TICKET - to be inserted] marks where the system-generated ticket and its cited evidence will be embedded for each case.

- Clear cases -

Case 1 (H01 - Policy/Compliance - expected: KG augmentation helps). *Trigger query:* "What are the compliance requirements for patching the Wi-Fi subsystem?" *Gold cross-source links:* CISA KEV (HTML + JSON), Samsung, Pixel bulletins. *Why included:* the answer lives in vocabulary-mismatched sources reachable via a shared CVE; this is the regime the knowledge graph is designed for. [ACTION TICKET - to be inserted]

Case 2 (B03 - Identifier-Focused - expected: clean, RAG-saturated). *Trigger query:* CVE-2025-0097 / CVE-2025-0096 / CVE-2025-0093 identifier extract. *Gold links:* Android Security Bulletin, CISA KEV JSON, Samsung, Pixel. *Why included:* high-salience identifier strings that the system links across sources with full coverage - a clean, unambiguous ticket. [ACTION TICKET - to be inserted]

Case 3 (D04 - Natural Language - expected: well-grounded synthesis). *Trigger query:* "Which Android security bulletins correspond to the CVEs added to the CISA Known Exploited Vulnerabilities catalog this month?" *Gold links:* Bulletin, CISA KEV (HTML), Samsung, Pixel. *Why included:* analyst-style question requiring multi-source synthesis; tests whether the ticket reads as a coherent, sourced answer. [ACTION TICKET - to be inserted]

Case 4 (B07 - Identifier-Focused, narrow scope - expected: correctly scoped). *Trigger query:* CVE-2025-0098 (Camera HAL) / CVE-2025-0091 (ActivityManagerService). *Gold links:* Bulletin, Samsung only. *Why included:* a case where the *correct* answer is a small link set - tests whether the ticket avoids over-linking and keeps priority proportionate. [ACTION TICKET - to be inserted]

- Ambiguous / edge cases -

Case 5 (C09 - Component-Focused - KG blackout, graceful degradation). *Trigger query:* "Knox attestation bypass allowed modified devices to pass integrity checks. Secure Folder race condition... active exploitation detected..." *Gold links:* Bulletin, Pixel. *Why included:* the regex extractor captures no entity for "Knox," so the graph contributes nothing and the system falls back to vector retrieval. Probes whether the ticket is still relevant/actionable despite the structural blackout, and whether urgency is right given "active exploitation." [ACTION TICKET - to be inserted]

Case 6 (A09 - Full Diff - fusion misfire, honest negative). *Trigger:* the raw Samsung bulletin diff (entity-rich). *Gold links:* Bulletin, CISA (HTML+JSON), NVD, Pixel. *Why included:* fusion pushes a true-positive OEM source past rank 5 - the case where KG augmentation *hurts* retrieval. Probes whether experts notice a missing/weakened evidence link and whether the ticket's confidence is still warranted. [ACTION TICKET - to be inserted]

Case 7 (G01 - Temporal-Change - extraction-induced recall collapse). *Trigger query:* "What new kernel vulnerabilities were added between the original and updated March bulletin?" *Gold links:* CISA (HTML+JSON), Pixel, Samsung. *Why included:* temporal "what changed" phrasing yields no extractable identifier, the KG-only path collapses, and recall is fragile. Probes evidence sufficiency and whether the ticket over- or under-states the delta. [ACTION TICKET - to be inserted]

Case 8 (J03 - Noisy/Adversarial - robustness + calibration). *Trigger query:* "If I were a hacker, how would I use the ARM Mali GPU vulnerability?" *Gold links:* Bulletin, CISA KEV JSON, Samsung, Pixel. *Why included:* adversarially framed input around a real CVE. Probes whether the ticket stays on the legitimate monitoring task (relevance, refusal

of the misuse framing) and whether priority/urgency remains appropriate. [ACTION TICKET - to be inserted]

Per-Case Rating Sheet (repeat for Cases 1-8)

Case	D1 Relevance	D2 Actionability	D3 Evidence	D4 Priority/Urgency
1				
2				
3				
4				
5				
6				
7				
8				

Per-case comments (optional):

- Case 1: ...
 - Case 2: ...
 - Case 3: ...
 - Case 4: ...
 - Case 5: ...
 - Case 6: ...
 - Case 7: ...
 - Case 8: ...
-

Debrief (brief, free-text)

1. **Overall trust.** Would you rely on these tickets to triage real ecosystem changes? Why / why not?
 2. **Strongest / weakest dimension.** Across the eight cases, where did the system most and least meet your expectations?
 3. **Failure visibility.** In the ambiguous cases (5-8), was it apparent *to you* that something was off, or did the ticket read as confident regardless?
 4. **Missing dimensions.** Is there a quality aspect this rubric does not capture that you would want scored?
 5. **One change.** If you could change one thing about the tickets, what would it be?
-

Notes on Scope and Refinement

- **Draft status.** Case selection, rubric anchors, and debrief questions are an

initial proposal and will be refined with expert input, per the advisor's note.

- **Inter-rater reliability.** With multiple experts we can report agreement (e.g., Krippendorff's alpha or mean pairwise weighted kappa) per dimension; the 8-case size is a pilot, expandable if reliability is low.
- **Provenance.** Cases map directly to the manuscript's failure analysis (Section 7.3) and case studies (Section 7.4), so expert scores can be discussed against the quantitative retrieval results.
- **Open item.** The generated Action Tickets for each case are pending; they will be produced from the pipeline and embedded at the [ACTION TICKET] markers before the experts review.